

Hardware Trojan Detection Using Graph Neural Networks

AST-based Analysis with GCN Architecture

Md Omar Faruque

May 8, 2025

Outline

1 Introduction

2 AST Generation

3 GCN Architecture

4 Experimental Results

Hardware Trojans: A Significant Threat

- **Hardware Trojans:** Malicious modifications to circuit designs
- Often inserted during third-party manufacturing or through IP cores
- Can lead to:
 - Information leakage
 - Denial of service
 - Functional changes
- **Challenge:** Detection is difficult at RTL level
- **Our approach:** Use Graph Neural Networks on code representations

Verilog RTL & Graph Representations

Abstract Syntax Tree (AST)

- Represents code structure
- Language constructs as nodes
- Parent-child relationships
- Captures programming patterns

Our approach: Using AST representation for more effective trojan detection


What is an Abstract Syntax Tree (AST)?

Simple definition:

- AST is a tree representation of code's structure
- Every line of code becomes a branch in the tree
- Similar to how sentences have grammar structure

Advantages for trojan detection:

- Preserves programming patterns
- Shows suspicious control flow
- Captures language-specific features



Simple AST for:

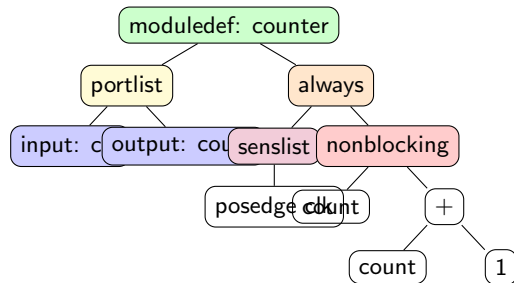
```
if (x == 0) {y = 1;}
```

AST Generation: A Simple Example

Original Verilog Code:

```
module counter(  
    input clk,  
    output reg [3:0]  
    count  
);  
    always @(posedge  
clk) begin  
        count <= count  
+ 1;  
    end  
endmodule
```

Simplified AST:



AST Node Types Discovered

- Our system discovered many Verilog constructs:
- Common node types:
 - moduledef (module definitions)
 - always (always blocks)
 - senslist (sensitivity lists)
 - if/case (conditionals)
 - blocking/nonblocking (assignments)
 - Found 45+ unique node types across all samples
 - rvalue/lvalue (values)
 - pointer (variable references)
 - eq/and/xor (operators)
 - partselect (bit selection)
 - concat (concatenation)

Sample Type Dictionary from Processing

- 0: source
- 1: description
- 2: moduledef
- 3: paramlist
- 4: portlist
- 5: decl
- 6: instancelist
- 7: ioport
- 8: input
- 9: width
- 10: intconst
- 19: block
- 20: sens
- 21: if
- 22: eq
- 23: nonblocking
- 24: lvalue
- 25: rvalue
- 26: pointer
- 27: xor
- 28: land

AST Feature Extraction

Our system creates rich node features from AST:

Process: Verilog Code → Parse into AST → Convert to Graph → Extract Features


- **Type encoding:** One-hot encoding of node type (50 dimensions)
- **Structural features:**
 - Depth in AST (hierarchical position)
 - Number of children (node complexity)
 - Average child type ID (contextual information)
- **Edge information:** Parent-child relationships
- These features capture both:
 - Local code structure
 - Global program patterns

Training Features: How We Represent AST Nodes

Example Node Feature Vector:

Feature Type	Value
Type: moduledef	1 (one-hot: [0,0,1,0,...])
Depth in AST	2
Number of children	5
Average child type ID	8.2

One-hot encoding example:



Why these features work:

- Type encoding captures node function
- Depth shows position in hierarchy
- Number of children indicates complexity
- Average child type provides context

Full feature vector size:

- 50 dimensions for type (one-hot)
- 3 dimensions for structural info


What is a Graph Convolutional Network (GCN)?

Simple definition:

- Neural network designed for graph data
- Learns patterns in connected data
- Like CNNs for images, but for graphs


Main operations:

- Message passing between nodes
- Neighborhood aggregation
- Feature transformation



Node A updates by receiving messages from neighbors B and C

GCN: A Working Example



Simple graph:
Initial features:

- Node A: $[1, 0]$
- Node B: $[0, 1]$
- Node C: $[1, 1]$
- Node D: $[0, 0]$

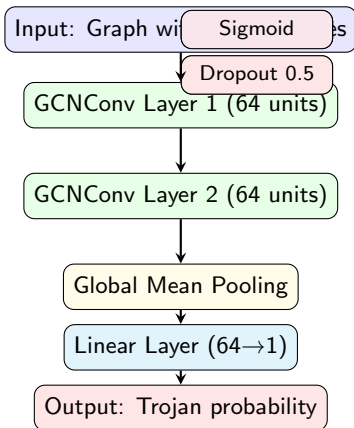
GCN layer computation:

- 1 Each node collects neighbor features
- 2 Apply weight matrix transformation
- 3 Add non-linearity (ReLU)

Example for Node A:

- Initial feature: $[1, 0]$
- Neighbors: B, C
- Collect: $[1, 0] + [0, 1] + [1, 1]$
- Average: $[0.67, 0.67]$
- Apply weights: $W \cdot [0.67, 0.67]^T$
- Apply ReLU: $\max(0, \text{result})$
- New feature: $[0.8, 1.2]$

GNN Implementation: Complete Architecture



GCNConv Layer Operations:

- 1 Calculate normalized adjacency matrix
- 2 Aggregate features from neighbors
- 3 Apply weight matrix transformation
- 4 Add non-linearity (ReLU)

Dropout (0.5): Randomly zeroes 50% of features during training to prevent overfitting

Global Mean Pooling: Averages all node features to get a single graph representation

Loss Function: Binary Cross Entropy with class weighting (0.26) to handle imbalance

Dataset Statistics

- **Dataset:** TJ-RTL-toy benchmark
- **Total circuits:** 43
 - 34 Trojan-infected (79%)
 - 9 Trojan-free (21%)
- **Class imbalance:** 0.26:1 (Clean:Trojan)
- **Circuits processed:**
 - AES, RC5, RC6, PIC16F84, RS232, VGA, XTEA, etc.
- **Node types discovered:** 45+

Experimental Setup

- **Cross-validation:** 2-fold stratified
- **Oversampling:** Random oversampling of minority class
- **Training:**
 - 100 epochs
 - Adam optimizer ($\text{lr}=0.001$)
 - BCELoss with class weighting
 - Batch size: 32
- **Model:** GCN with 2 layers (64 hidden units)
- **Random seed:** 42 (for reproducibility)

Implementation Details

Key Libraries:

- **PyTorch** (v1.10+)
- **PyTorch Geometric** (PyG)
- **PyVerilog** parser
- **NetworkX** for graphs
- **NumPy/Scikit-learn**
- **Matplotlib**

GNN Components:

- **GCNConv** (primary)
- **global_mean_pool**

Hardware/Runtime:

- **GPU:** NVIDIA RTX 3060 mobile
- **Training Time:** 30s
- **Codebase:** Python 3.9

Performance Metrics

Metric	Fold 1	Fold 2	Average
Accuracy	0.5909	0.9048	0.7478
Precision	0.9000	1.0000	0.9500
Recall	0.5294	0.8824	0.7059
F1 Score	0.6667	0.9375	0.8021

Table 1: Performance metrics across folds

Average Confusion Matrix

		Predicted	
		Clean	Trojan
2*Actual	Clean	4.0	0.5
	Trojan	5.0	12.0

Key Findings

- **High precision (95%):** Few false positives
- **Good F1 score (80.2%):** Balanced performance
- **Improvement across folds:** Fold 2 significantly better

Discussion & Limitations

■ Strengths:

- High precision - few false positives
- Effective at capturing code-level trojans

■ Limitations:

- Small dataset (43 circuits)
- Class imbalance (more trojans than clean)
- Fold variation indicates potential overfitting

■ Future work: Larger datasets, more complex architectures

Conclusions: What We've Learned

Key achievements:

- Used code structure (AST)
- Achieved 75% accuracy, 95% precision

Next steps:

- Larger, more diverse datasets
- Hyperparameter optimization
- Hybrid architecture (GNN+MLP)
- Neural architecture search

Thank You!

Questions?